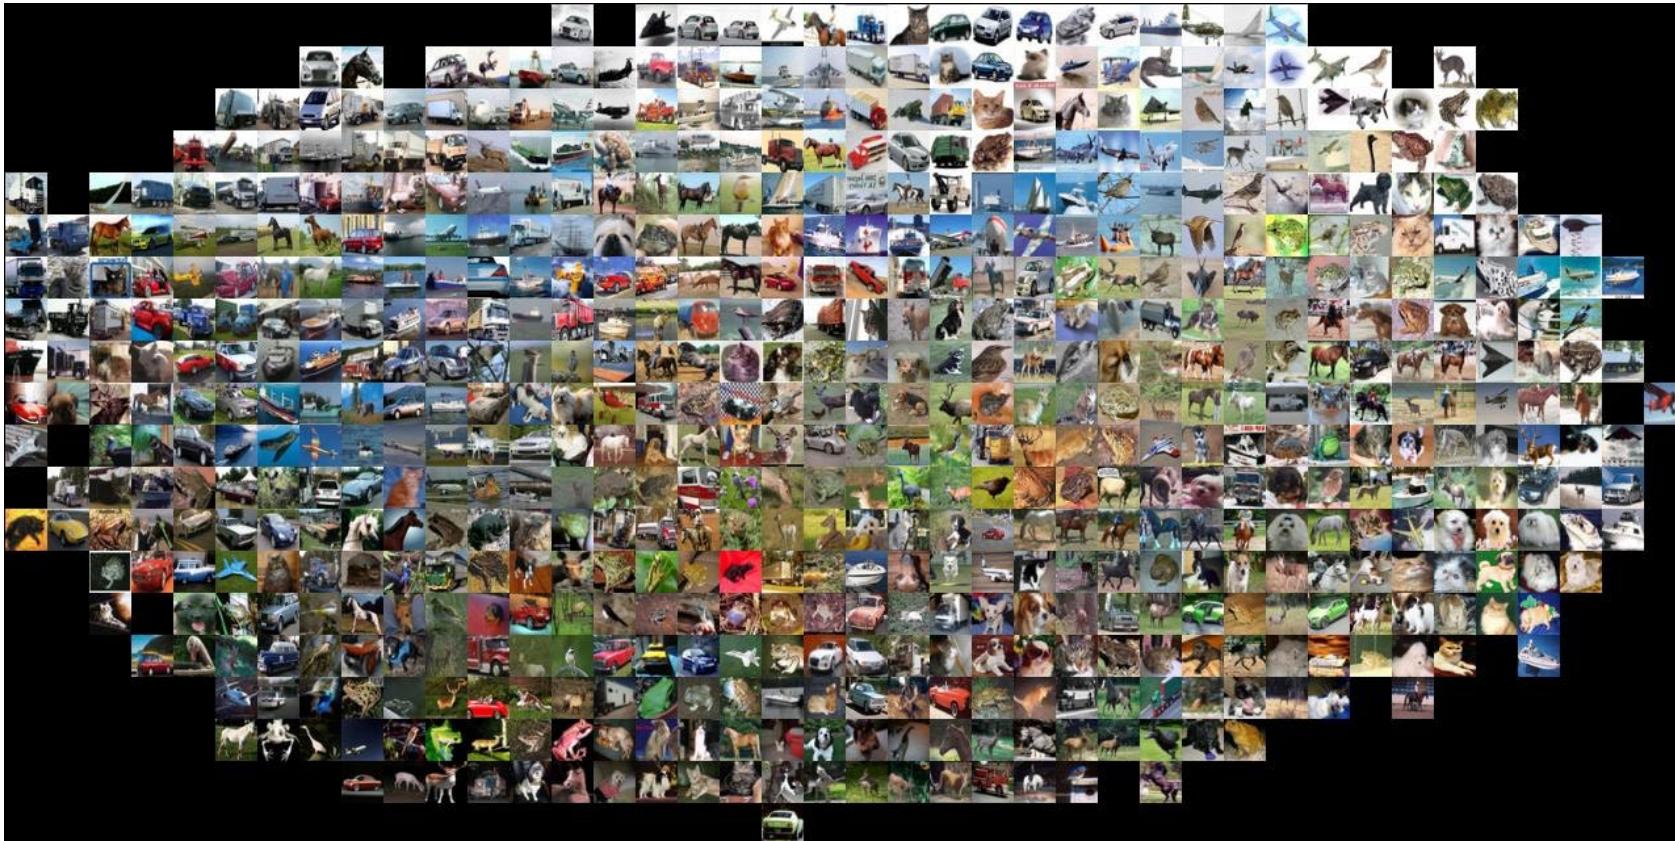


CS7GV1: Computer Vision

Image Classification



Credits: Some Slides from Noah Snavely and Abe Davis (Cornell)
and Fei-Fei Li, Justin Johnson, Serena Yeung (Stanford)
<http://vision.stanford.edu/teaching/cs231n/>

Image Classifiers in a Nutshell

- Input: an image
- Output: the class label for that image
- Label is generally one or more of the discrete labels used in training
 - e.g. {cat, dog, cow, toaster, apple, tomato, truck, ... }

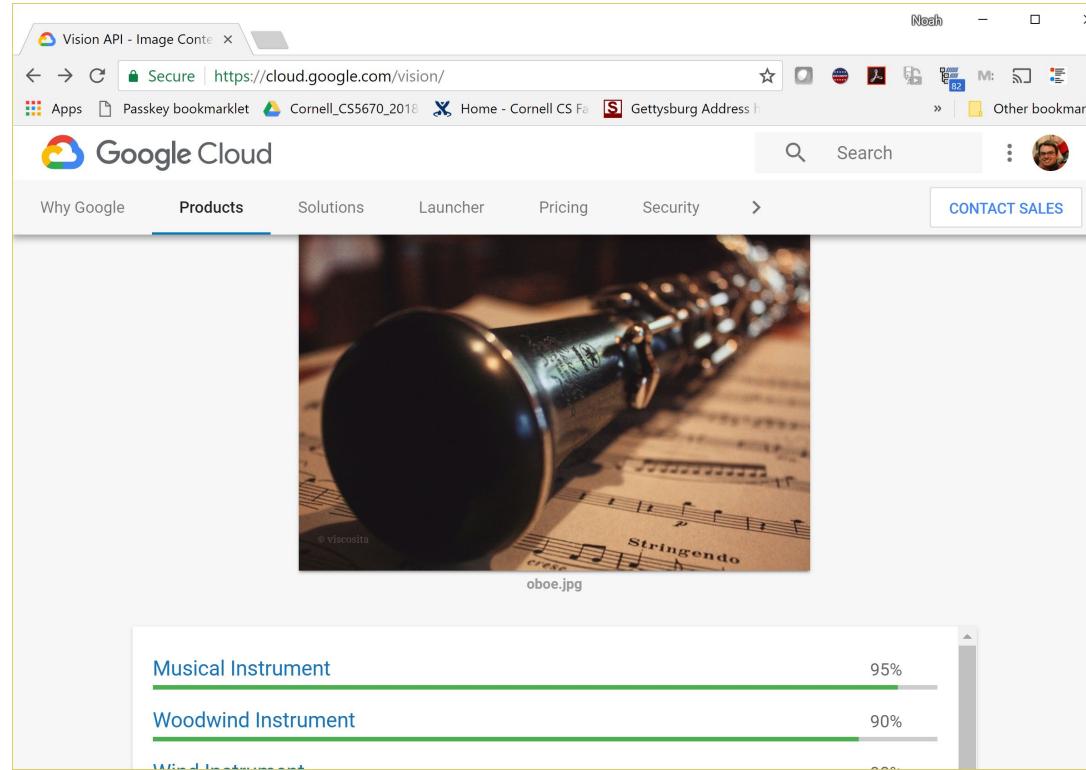
```
def classifier(image):  
    //Do some stuff  
    return class_label;
```

$$f\left(\begin{array}{c} \text{Cat Image} \end{array}\right) = \text{"Cat"}$$

$$f\left(\begin{array}{c} \text{Dog Image} \end{array}\right) = \text{"Dog"}$$

$$f\left(\begin{array}{c} \text{Toaster Image} \end{array}\right) = \text{"Toaster"}$$

Image classification demo



<https://cloud.google.com/vision/docs/drag-and-drop>

See also:

<https://aws.amazon.com/rekognition/>

<https://www.clarifai.com/>

<https://azure.microsoft.com/en-us/services/cognitive-services/computer-vision/>

The Semantic Gap



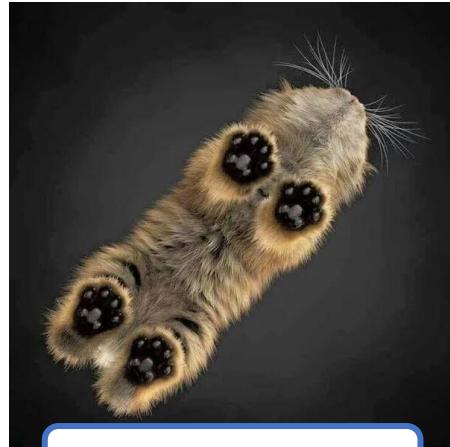
What we see

```
01110 111010010111  
11010111 1101010101  
0 01001 1111010111  
101010111101110111  
11001111001 1 01 10  
010111 1111101101  
01 11101 10010010  
10011111 100011100  
10100100100110001  
10000 100011101110  
011000001111 1 10  
1 1011001 010011 10  
11 10001 1111 101  
111 101010 10111 :  
100010101001010110  
1001 1111000010 110  
10111100001111 110:
```

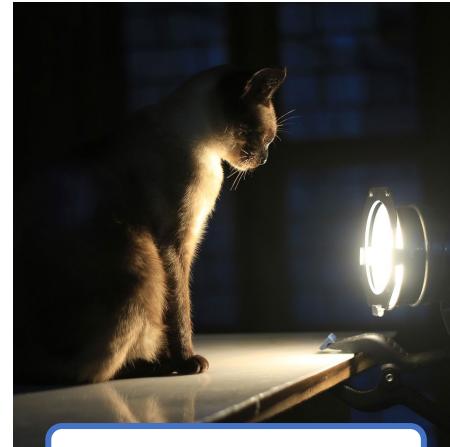
What the computer sees

Variation Makes Recognition Hard

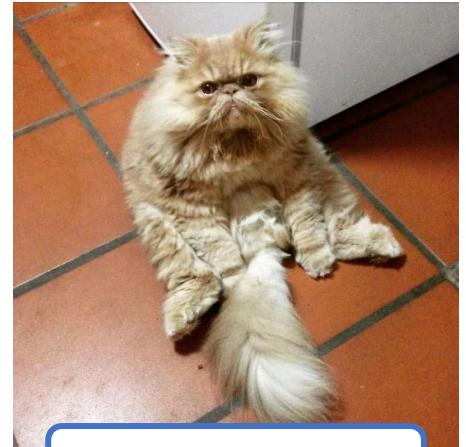
- The same class of object can appear *very* differently in different images



Viewpoint Variation



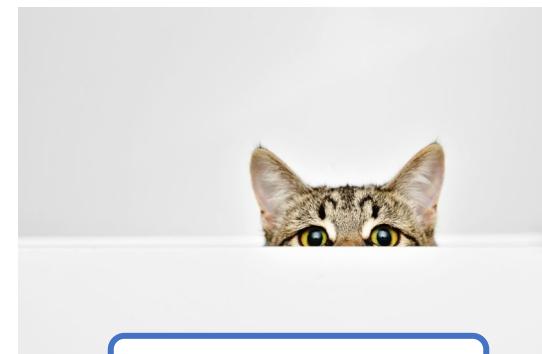
Lighting Variation



Deformation



Background Clutter



Occlusion

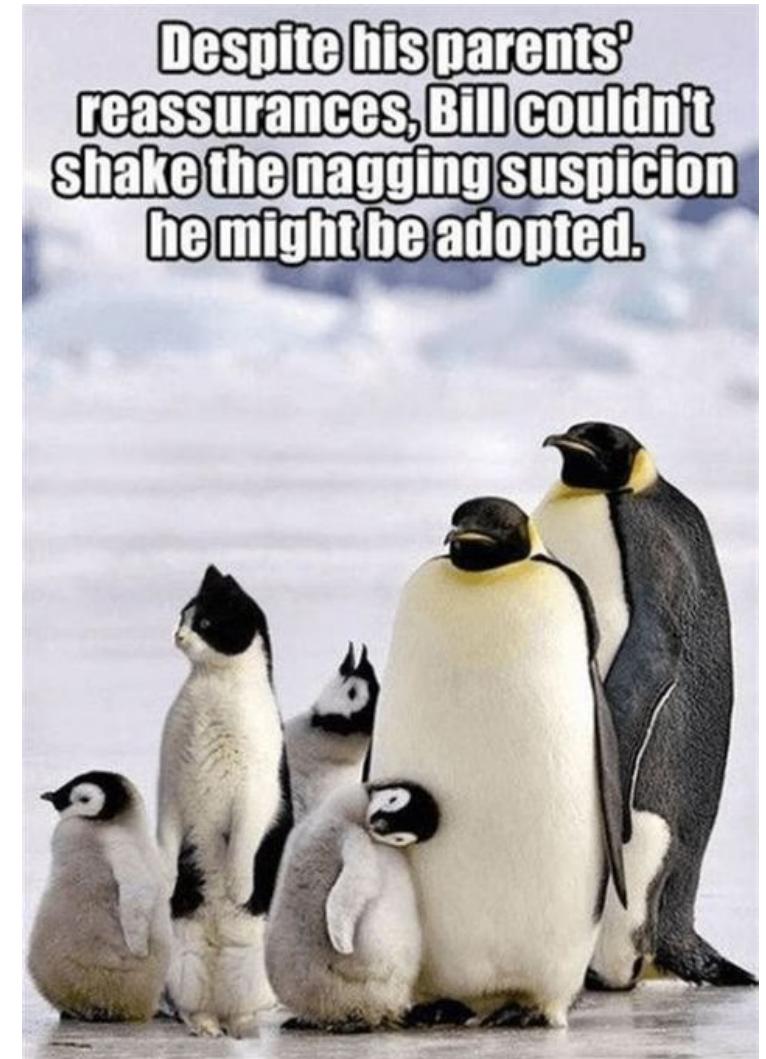
The Problem is Under-constrained

- Distinct realities can produce the same image...



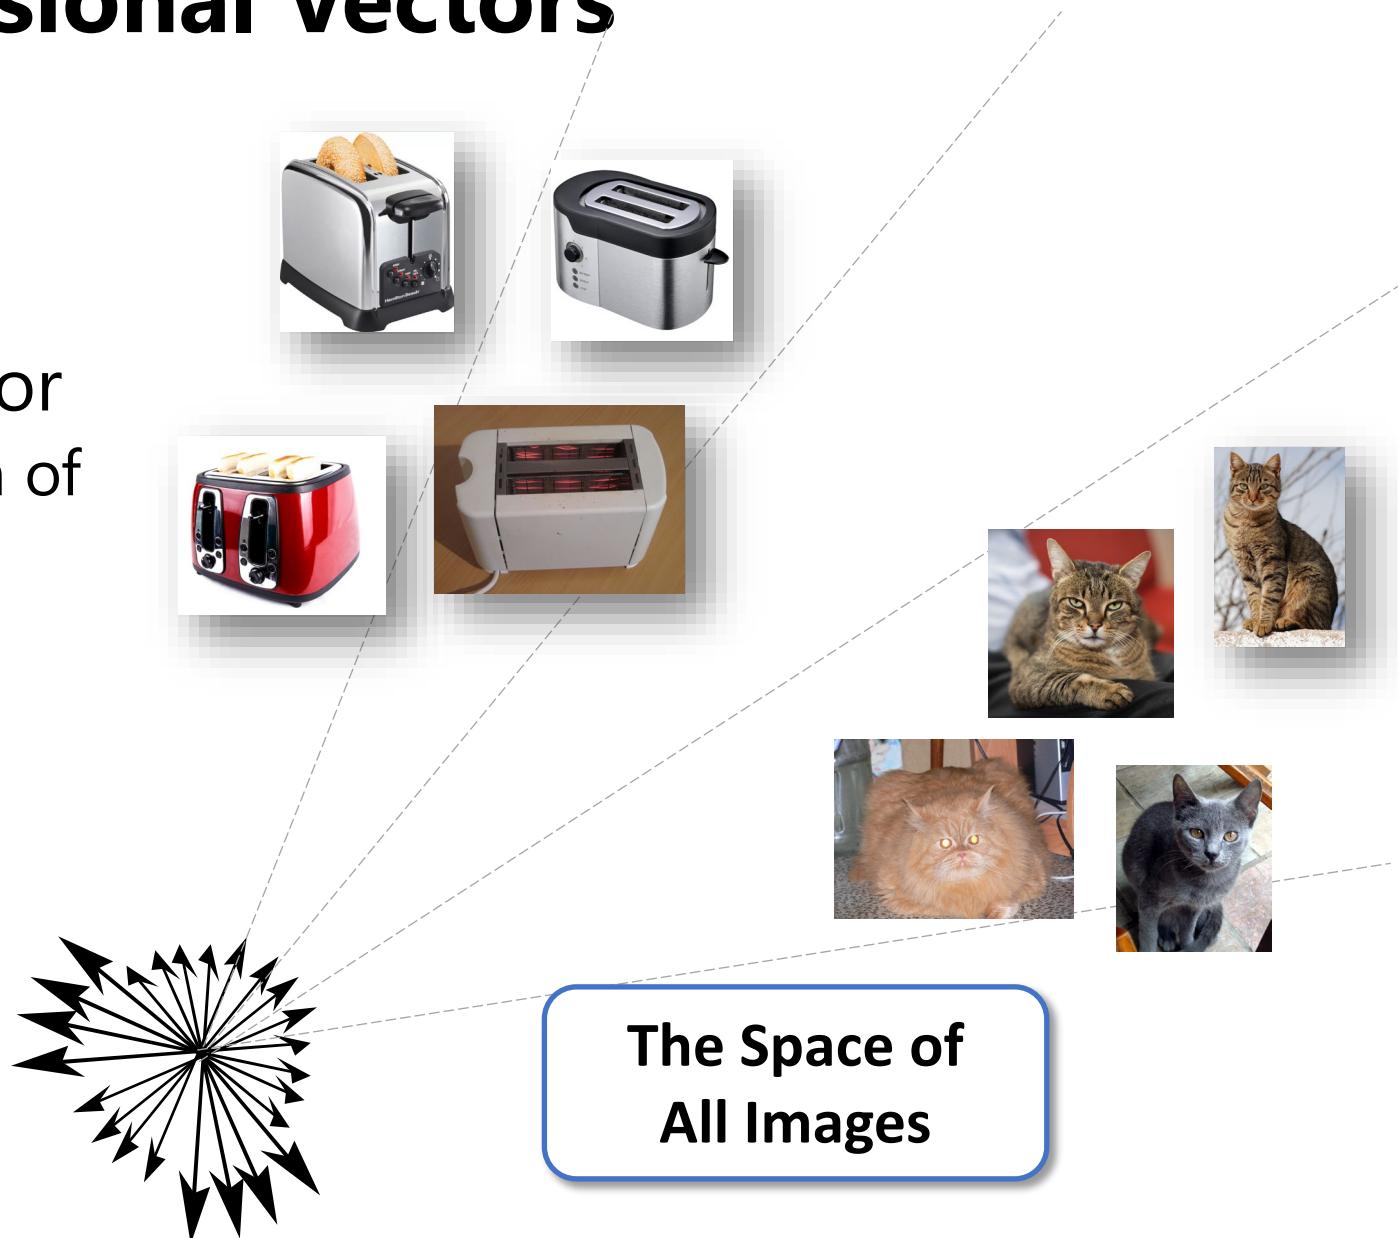
The Problem is Under-constrained

- Distinct realities can produce the same image...
- We generally can't compute the "right" answer, but we can compute the most likely one...
- We need some kind of prior to condition on. We can learn this prior from data.



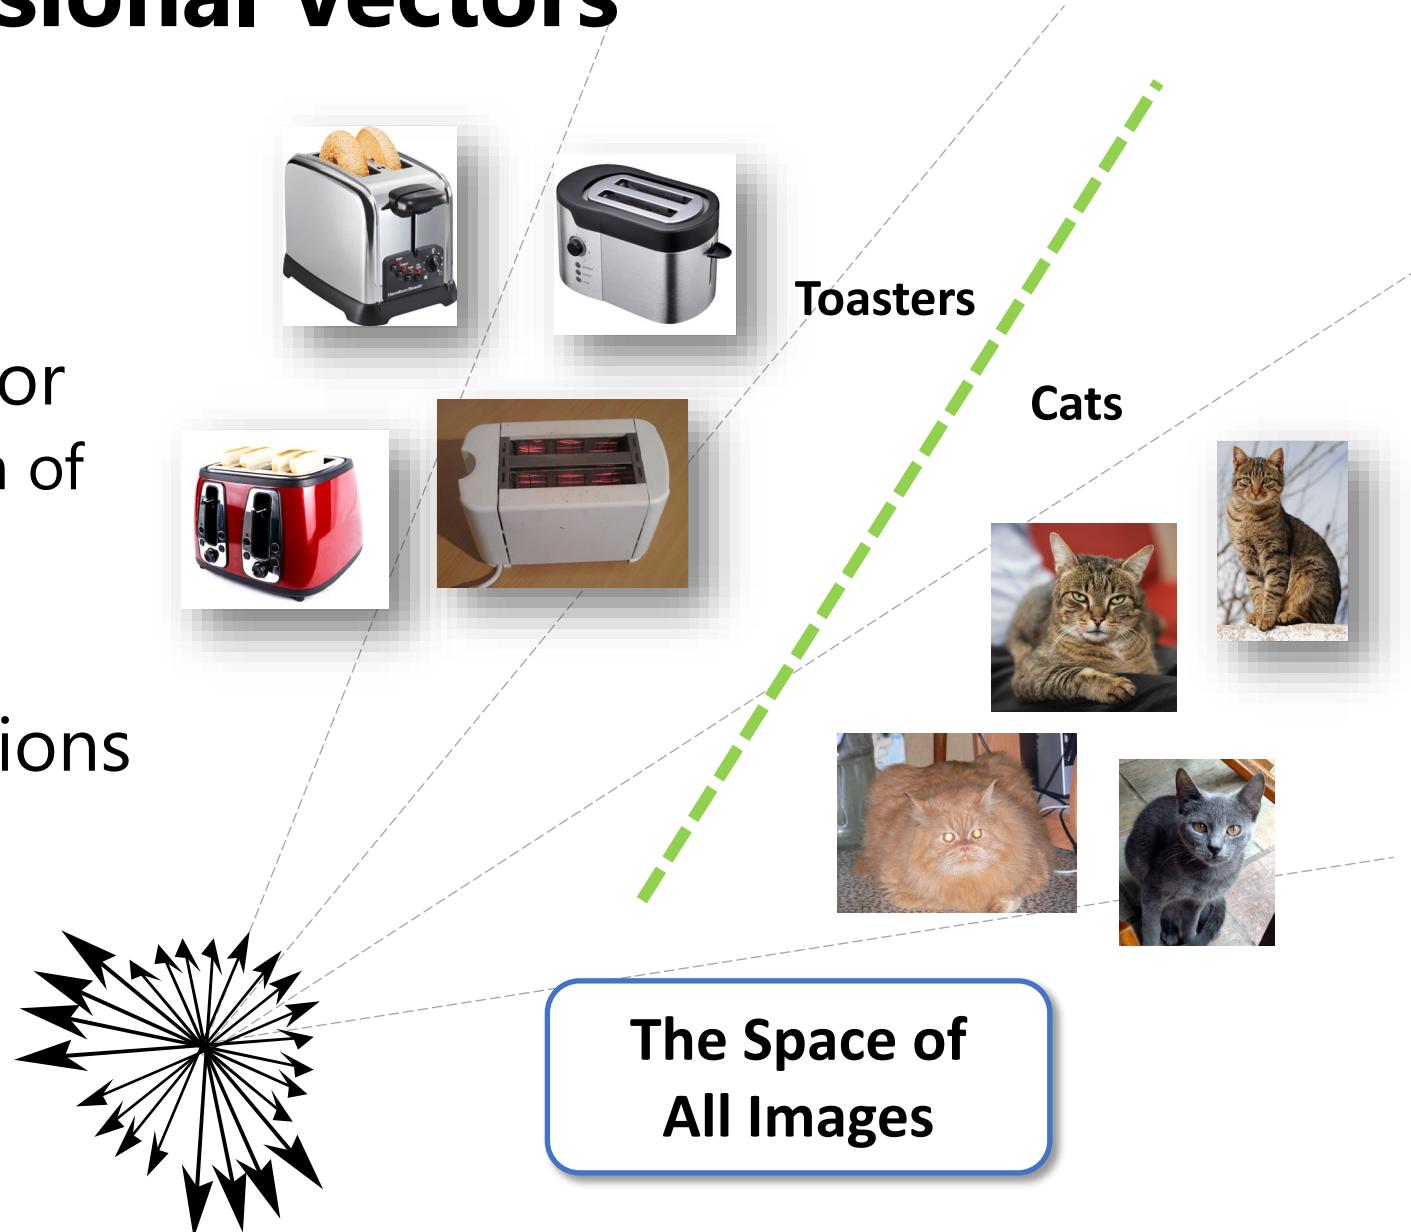
Images As High-Dimensional Vectors

- An image is just a bunch of numbers
- Let's stack them up into a vector
 - Our training data is just a bunch of high-dimensional points now



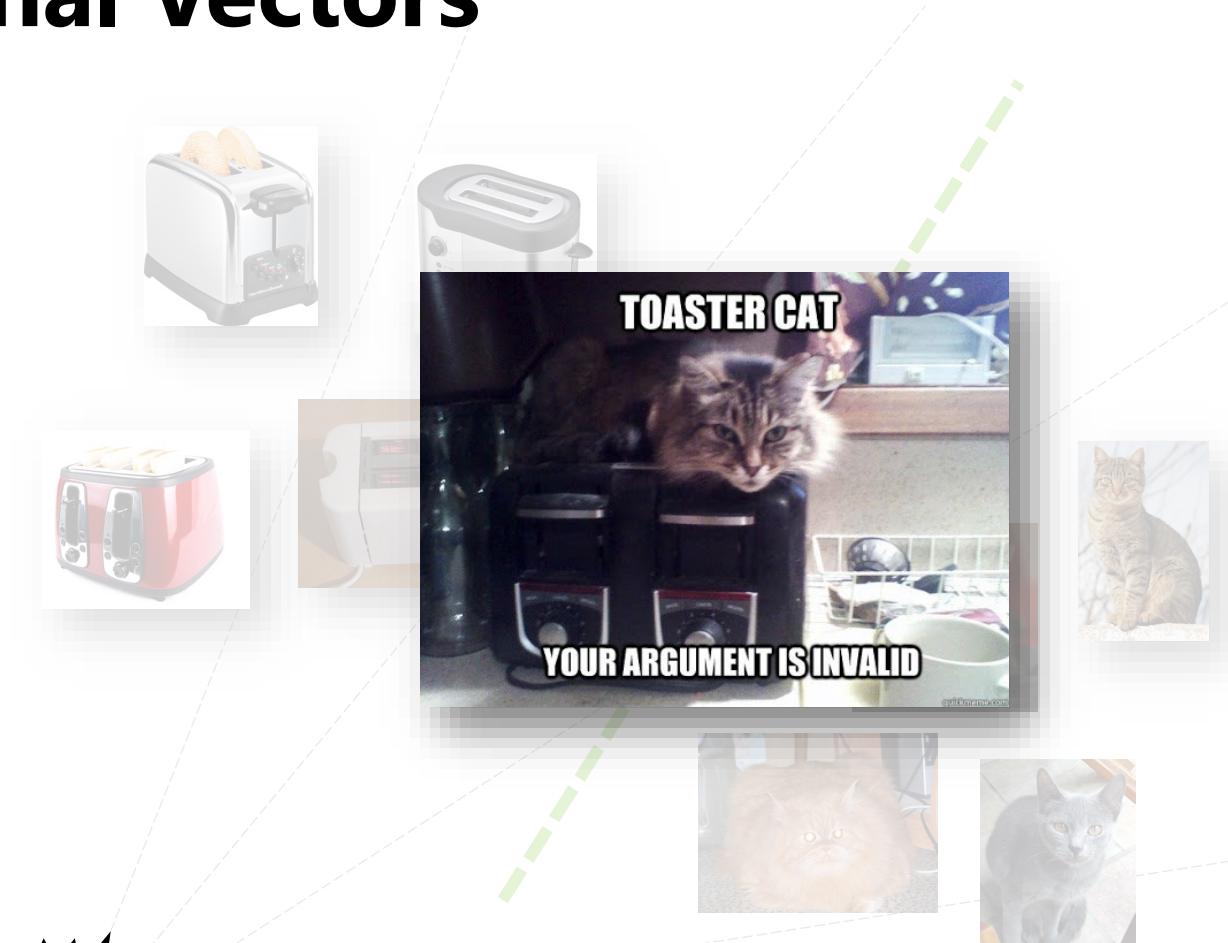
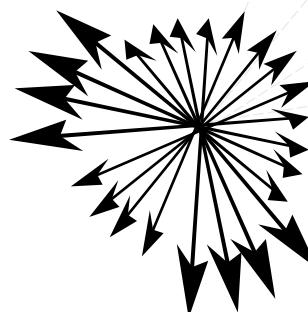
Images As High-Dimensional Vectors

- An image is just a bunch of numbers
- Let's stack them up into a vector
 - Our training data is just a bunch of high-dimensional points now
- Divide space into different regions for different classes



Images As High-Dimensional Vectors

- An image is just a bunch of numbers
- Let's stack them up into a vector
 - Our training data is just a bunch of high-dimensional points now
- Divide space into different regions for different classes



The Space of
All Images

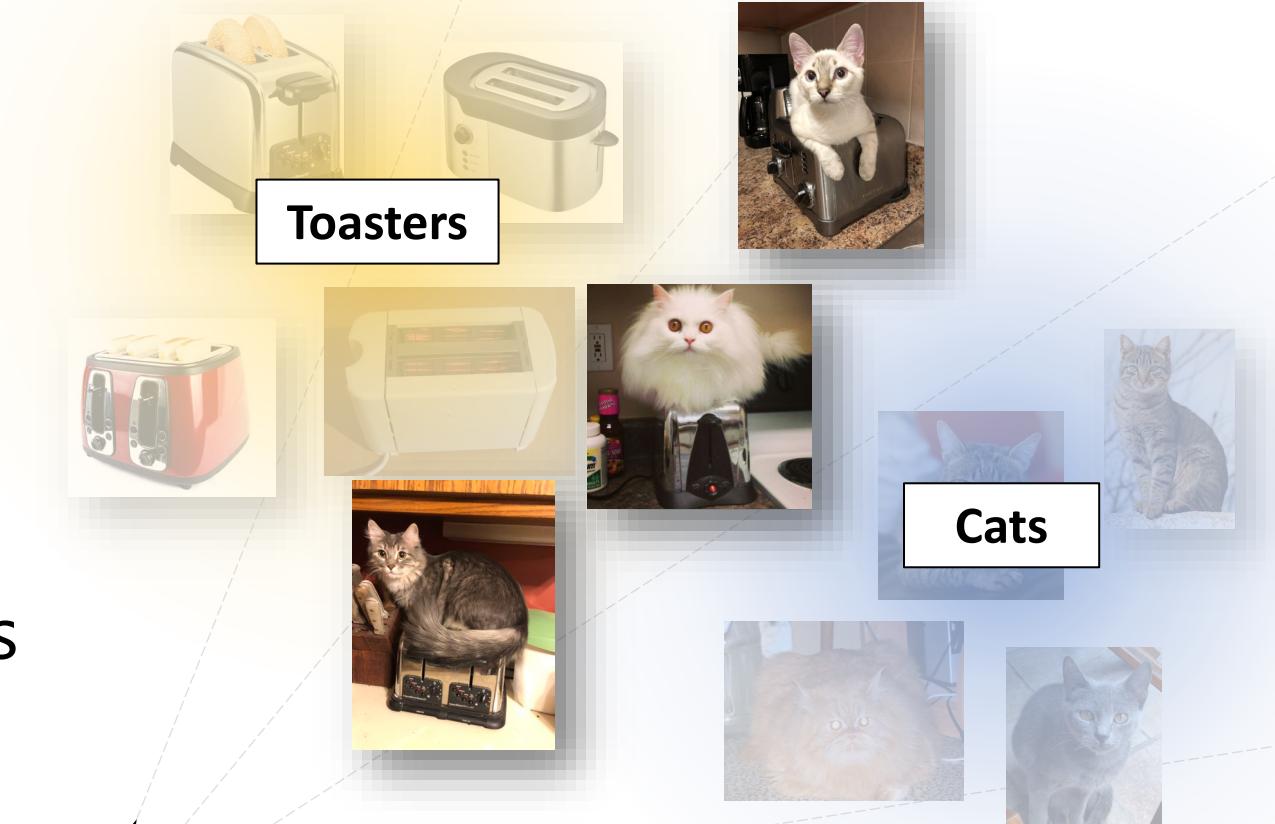
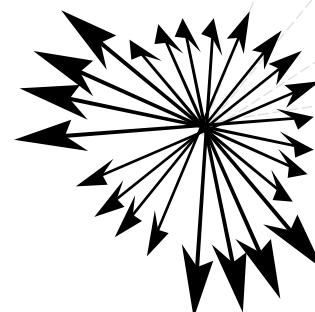
Images As High-Dimensional Vectors

- An image is just a bunch of numbers
- Let's stack them up into a vector
 - Our training data is just a bunch of high-dimensional points now

- Divide space into different regions for different classes

or

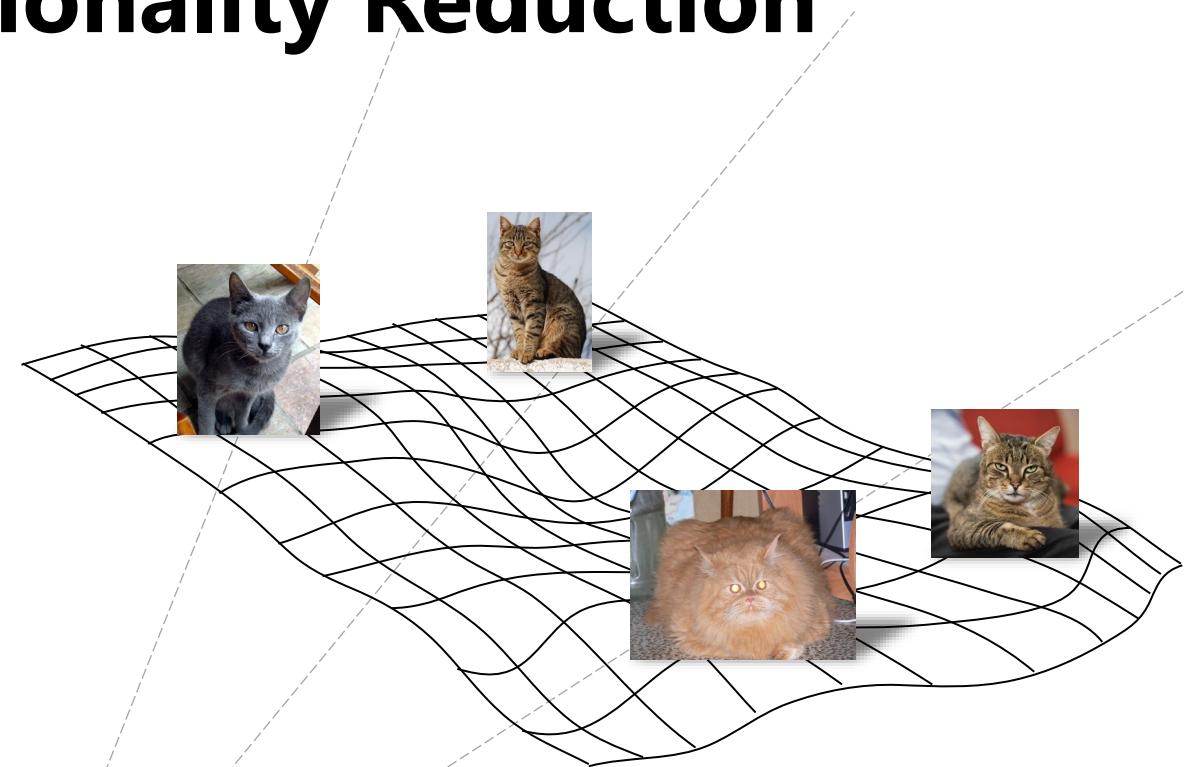
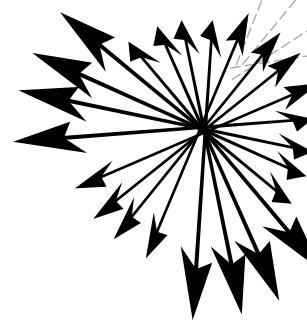
- Define a distribution over space for each class



The Space of
All Images

Image Features and Dimensionality Reduction

- How high-dimensional is an image?
 - Let's consider an iPhone X photo:
 - 4032 x 3024 pixels
 - Every pixel has 3 colors
 - 36,578,304 pixels (36.5 Mega pixels)
- In practice, images sit on a lower-dimensional manifold
- Think of image features and dimensionality reduction as ways to represent images by their location on such manifolds



The Space of
All Images

Training & Testing a Classifier

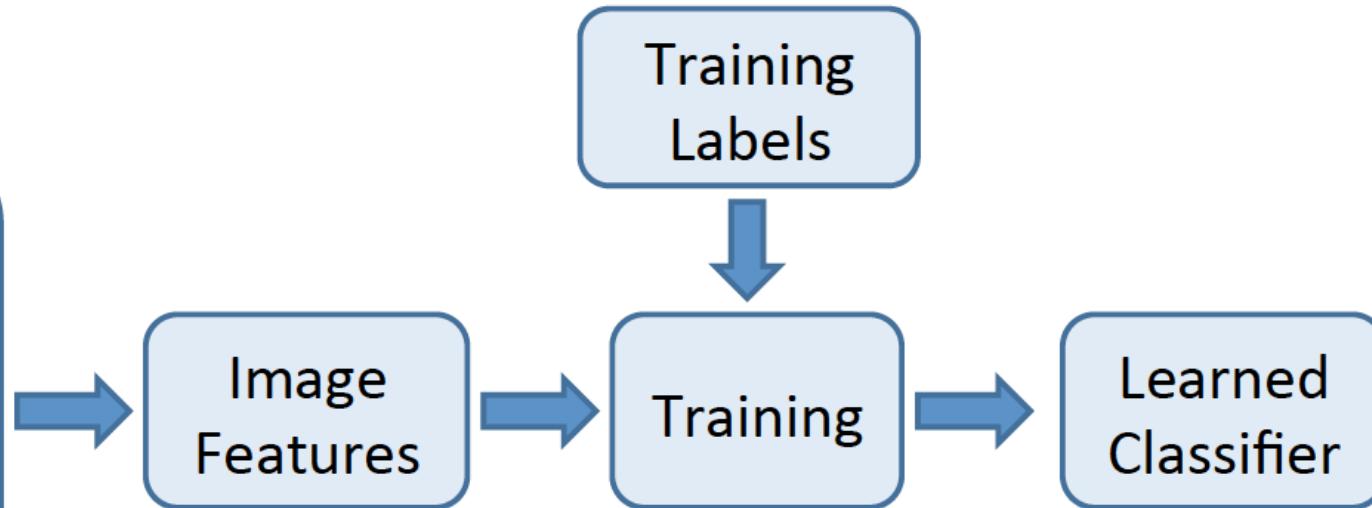
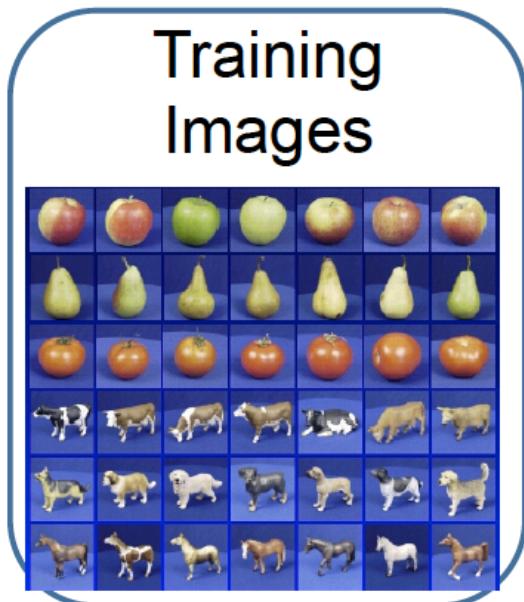
- Collect a database of images with labels
- Use ML to train an image classifier
- Evaluate the classifier on test images

Example training set



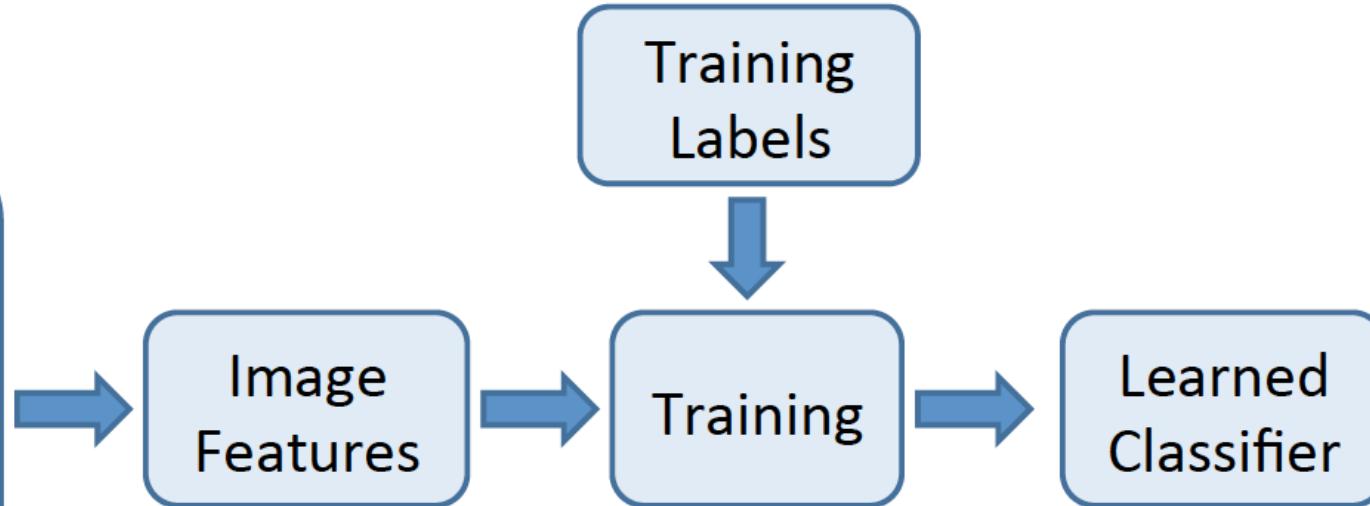
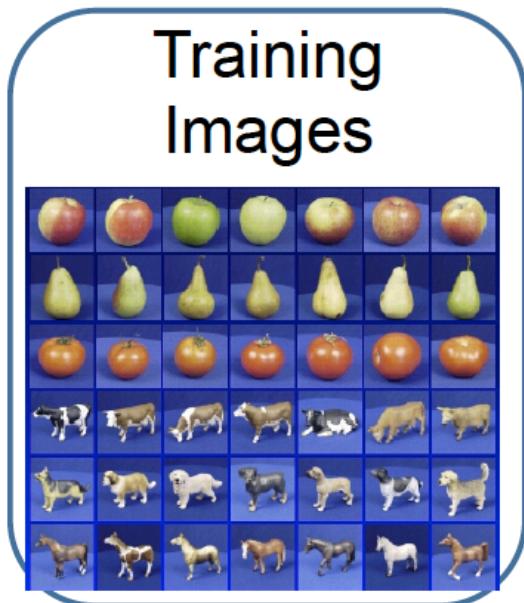
Training & Testing a Classifier

Training



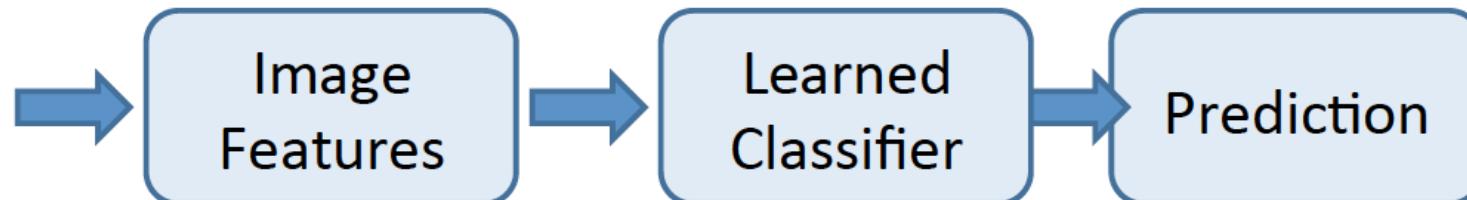
Training & Testing a Classifier

Training



Testing

Test Image

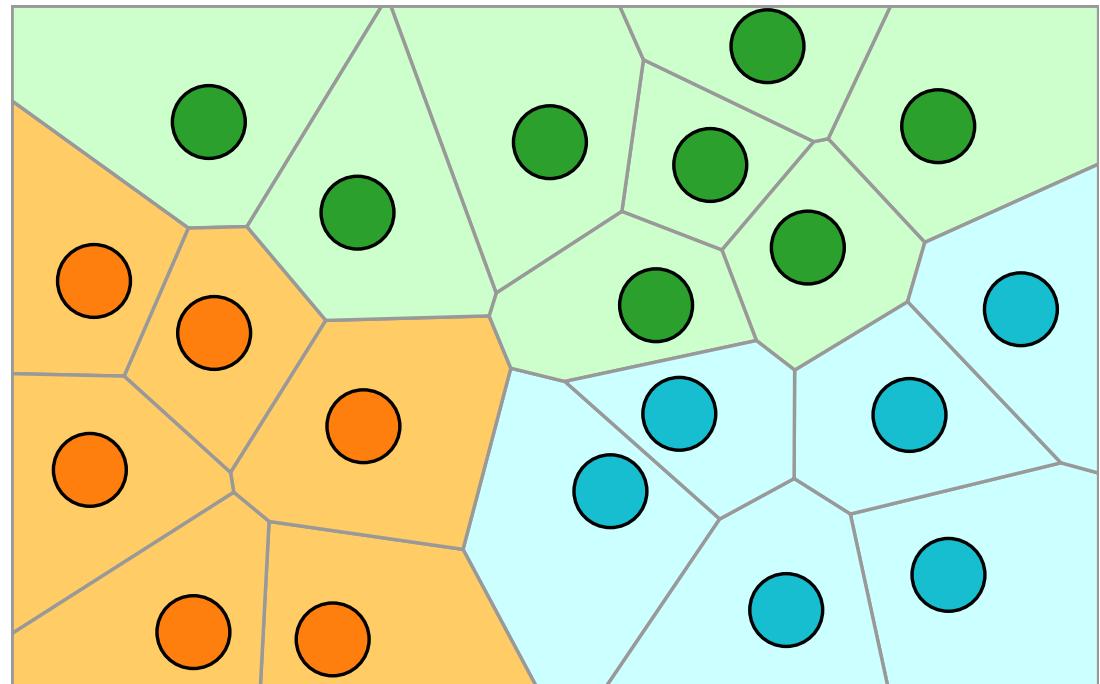


Classifiers

- Nearest Neighbor
- kNN (“k-Nearest Neighbors”)
- Linear Classifier
- Neural Network
- Deep Neural Network
- ...

First: Nearest Neighbor (NN) Classifier

- Train
 - Remember all training images and their labels
- Predict
 - Find the closest (most similar) training image
 - Predict its label as the true label



CIFAR-10 and NN results

Example dataset: **CIFAR-10**

10 labels

50,000 training images

10,000 test images.



CIFAR-10 and NN results

Example dataset: **CIFAR-10**

10 labels

50,000 training images

10,000 test images.

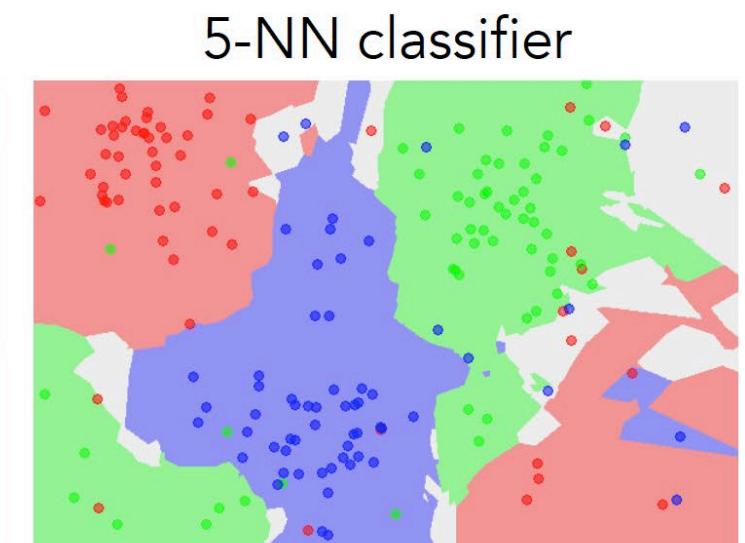
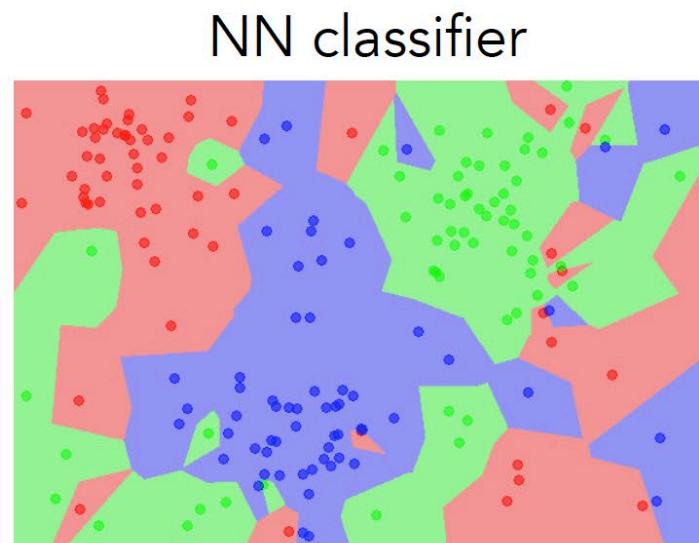
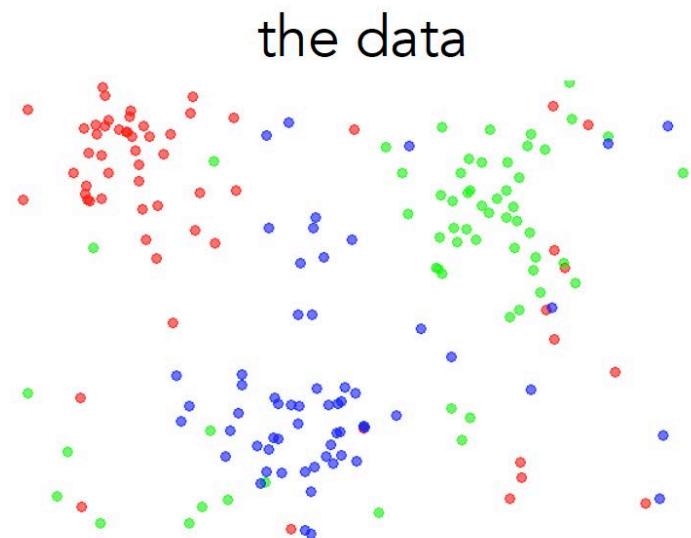


For every test image (first column),
examples of nearest neighbors in rows



k-nearest neighbor

- Find the k closest points from training data
- Take **majority vote** from K closest points



What does this look like?



What does this look like?



How to Define Distance Between Images

L1 distance:

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

Where I_1 denotes image 1,
and p denotes each pixel

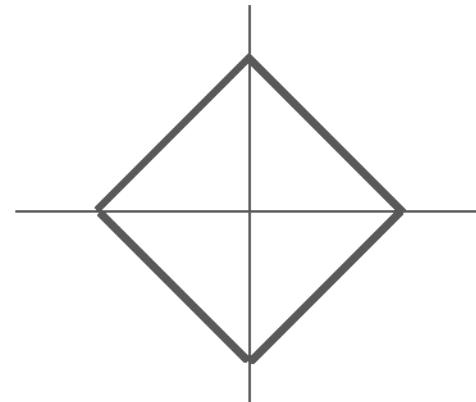
test image				training image				pixel-wise absolute value differences				→ 456
56	32	10	18	10	20	24	17	46	12	14	1	
90	23	128	133	8	10	89	100	82	13	39	33	
24	26	178	200	12	16	178	170	12	10	0	30	
2	0	255	220	4	32	233	112	2	32	22	108	

Choice of distance metric

- Hyperparameter

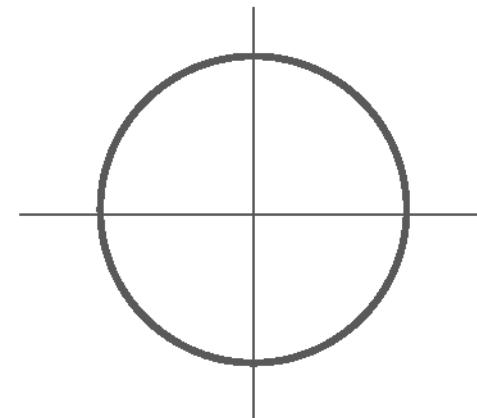
L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



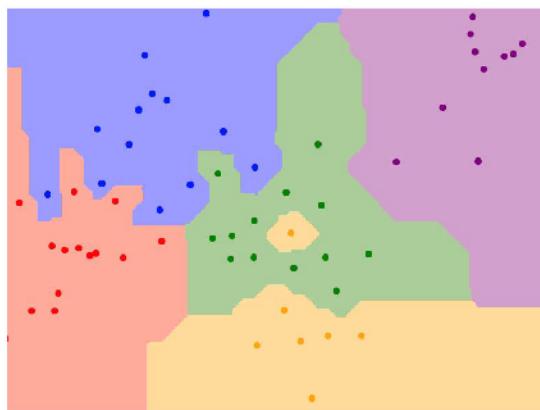
- Two most commonly used special cases of p-norm

$$\|x\|_p = \left(|x_1|^p + \cdots + |x_n|^p \right)^{\frac{1}{p}} \quad p \geq 1, x \in \mathbb{R}^n$$

K-Nearest Neighbors: Distance Metric

L1 (Manhattan) distance

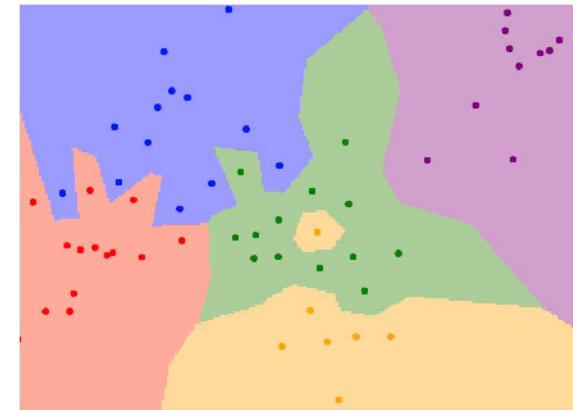
$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



$K = 1$

L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



$K = 1$

Demo: <http://vision.stanford.edu/teaching/cs231n-demos/knn/>

Hyperparameters

- What is the **best distance** to use?
- What is the **best value of k** to use?
- These are **hyperparameters**: choices about the algorithm that we set rather than learn
- How do we set them?
 - One option: try them all and see what works best

Setting Hyperparameters

Idea #1: Choose hyperparameters
that work best on the data

Your Dataset

Setting Hyperparameters

Idea #1: Choose hyperparameters
that work best on the data

BAD: $K = 1$ always works
perfectly on training data

Your Dataset

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

train

test

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

train

test

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

train

test

Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!

train

validation

test

Setting Hyperparameters

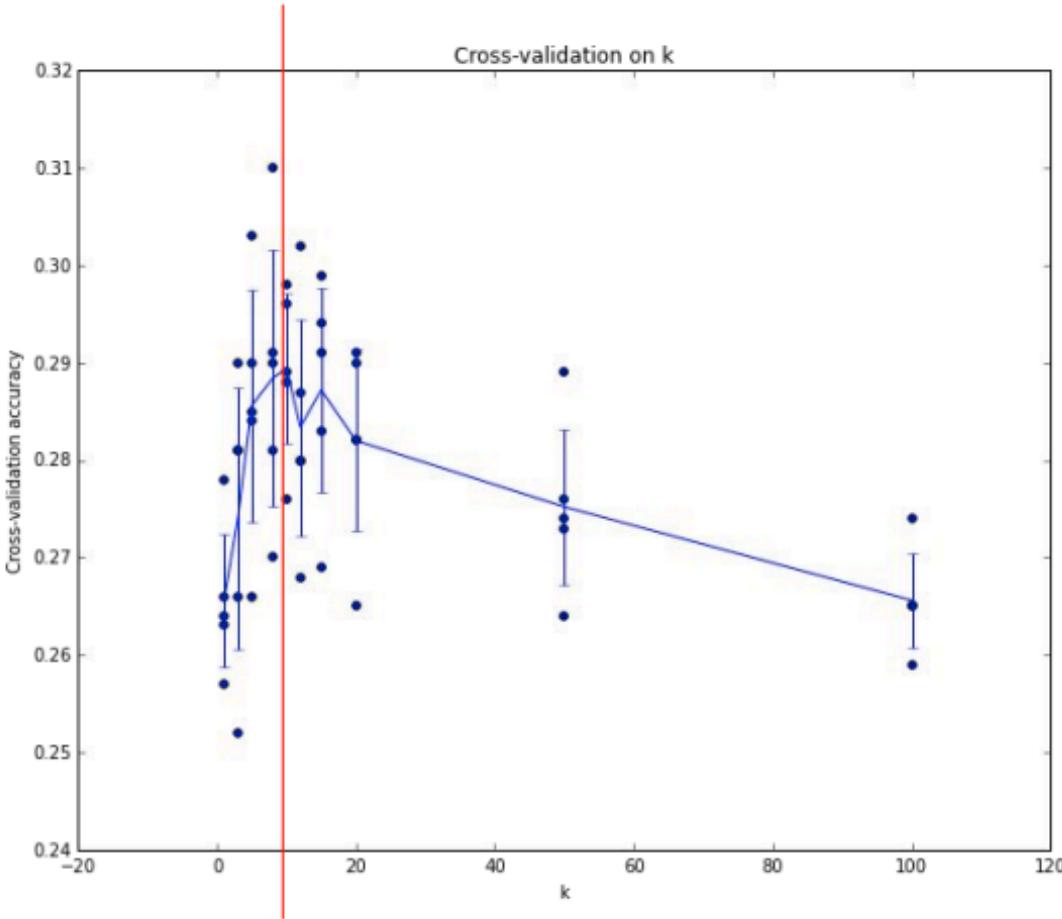
Your Dataset

Idea #4: Cross-Validation: Split data into **folds**,
try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

Hyperparameter Tuning



Example of
5-fold cross-validation
for the value of k .

Each point: single
outcome.

The line goes
through the mean, bars
indicated standard
deviation

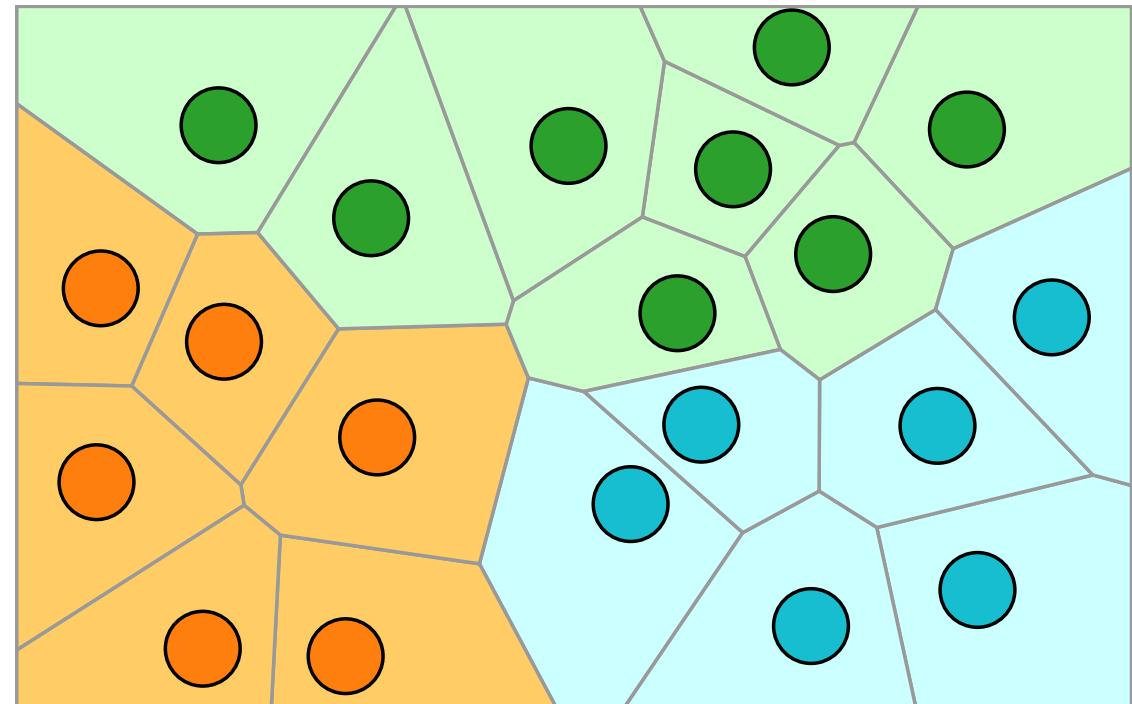
(Seems that $k \approx 7$ works best
for this data)

Recap: How to pick hyperparameters?

- Methodology
 - Train and test
 - Train, validate, test
- Train for original model
- Validate to find hyperparameters
- Test to understand generalizability

kNN -- Complexity and Storage

- N training images, M test images
- Training: $O(1)$
- Testing: $O(MN)$
- We often need the opposite:
 - Slow training is ok
 - Fast testing is necessary



k-Nearest Neighbors: Summary

- In **image classification** we start with a **training set** of images and labels, and must predict labels on the **test set**
- The **K-Nearest Neighbors** classifier predicts labels based on nearest training examples
- Distance metric and K are **hyperparameters**
- Choose hyperparameters using the **validation set**; only run on the test set once at the very end!

Problems with KNN: Distance Metrics

- terrible performance at test time
- distance metrics on level of whole images can be very unintuitive



(all 3 images have same L2 distance to the one on the left)

Problems with KNN: The Curse of Dimensionality

- As the number of dimensions increases, the same amount of data becomes more sparse.
- Amount of data we need ends up being exponential in the number of dimensions

